

Loop parallelization

Albert DeFusco

December 13, 2012

Compute π by Integration


- ▶ By definition

$$\pi = 4\arctan(1) \tag{1}$$

- ▶ Many formulas to compute π have centered around approximating $\arctan(1)$ by series
- ▶ $\arctan(1)$ is computed by integration


$$\pi = \int_0^1 \frac{4}{1+x^2} dx \tag{2}$$

Compute π by Integration



6 points : $\pi \approx 3.1124$

Compute π by Integration



41 points : $\pi \approx 3.1380$; 10^9 points : $\pi \approx 3.14159265$

Compute π in C++ and Fortran 90

```
1  #include <iostream>
2  using namespace std;
3
4  double arc(double x);
5  const long int num_steps=100000000;
6  double step;
7  int main()
8  {
9      double pi, sum=0.0;
10     step=1.0/(double) num_steps;
11
12     for (int i=0; i<=num_steps; i++)
13     {
14         double x=(i+0.5)*step;
15         sum+= arc(x);
16     }
17     pi = sum*step;
18
19     cout.precision(10);
20     cout << "pi is probably "
21         << fixed << pi << endl;
22
23     return 0;
24 }
25 double arc(double x)
26 {
27     double y = 4.0/(1+x*x);
28     return y;
29 }
```

```
1  program integratePi
2
3      implicit none
4      integer(kind=8) :: num_steps, i
5      real(kind=8) :: sum, step, pi, x
6
7      num_steps=100000000
8      sum=0.d0
9      step=1.d0/num_steps
10
11     do i=1, num_steps
12         x=(i+0.5d0)*step
13         sum=sum+arc(x)
14     enddo
15
16     pi=sum*step
17
18     write(6, &
19         '("pi is probably ",f12.10)') &
20         pi
21
22     contains
23     function arc(x)
24         implicit none
25         real(kind=8) :: arc, x
26         arc = 4.d0/(1.d0+x*x)
27     end function arc
28 end program integratePi
```

Compute π

```
>cp -r /home/sam/training/openmp/basic .  
>cd pi  
>module load gcc/4.5
```

```
#Compile in C++  
>CXX pi.cpp -o pi  
#Compile in Fortran 90  
>FC pi.f90 -o pi
```

```
>qsub pi-serial.job  
>cat pi.out  
pi is probably 3.1415926556
```

```
real    0m25.604s  
user    0m25.594s  
sys     0m0.000s
```

Variable Scopes

- ▶ Scope is the lexical context which determines the lifetime of a variable
- ▶ In C++ variables declared in the following regions are “locally” scoped
 - ▶ Regions in curly braces
 - ▶ Loops
 - ▶ Subroutines
- ▶ Fortran variables are valid within an entire subroutine and contained subroutines and functions
- ▶ Global variables exist in both languages
 - ▶ Variables declared as public in classes or outside of `main` in C++
 - ▶ Common blocks or modules in Fortran

Variable Scopes

- ▶ *lexical scope*
 - ▶ The physical text where a variable is valid
- ▶ *dynamical scope*
 - ▶ The runtime scope of a variable
- ▶ Scopes will be extremely important to OpenMP
 - ▶ Several clauses exist to control the scoping behavior

Variable Scopes

```
1  #include <iostream>
2  using namespace std;
3
4  double arc(double x);
5  const long int num_steps=100000000;
6  double step;
7  int main()
8  {
9      double pi, sum=0.0;
10     step=1.0/(double) num_steps;
11
12     for (int i=0; i<=num_steps; i++)
13     {
14         double x=(i+0.5)*step;
15         sum+= arc(x);
16     }
17     pi = sum*step;
18
19     cout.precision(10);
20     cout << "pi is probably "
21          << fixed << pi << endl;
22
23     return 0;
24 }
25 double arc(double x)
26 {
27     double y = 4.0/(1+x*x);
28     return y;
29 }
```

- ▶ `num_steps` and `step` are global
- ▶ `pi`, `step` and `sum` are valid for all of `main`
- ▶ `i` and `x` are valid only within the `for` loop

Variable Scopes

```
1  program integratePi
2
3  implicit none
4  integer(kind=8) :: num_steps, i
5  real(kind=8) :: sum, step, pi, x
6
7  num_steps=100000000
8  sum=0.d0
9  step=1.d0/num_steps
10
11 do i=1, num_steps
12     x=(i+0.5d0)*step
13     sum=sum+arc(x)
14 enddo
15
16 pi=sum*step
17
18 write(6, &
19     '("pi is probably ",f12.10)') &
20     pi
21
22 contains
23     function arc(x)
24         implicit none
25         real(kind=8) :: arc, x
26         arc = 4.d0/(1.d0+x*x)
27     end function arc
28 end program integratePi
```

- ▶ All variables declared in program are available to arc
- ▶ Variables in function arc are private to arc

Loop parallelization

- ▶ The `omp parallel for` region
 - ▶ Reads the loop bounds and divides the work
 - ▶ `i` becomes private to each thread in Fortran and C++
 - ▶ What other scopes do we expect?

```
1  #pragma omp parallel for
2  for (int i=0; i<=num_steps; i++)
3  {
4      double x=(i+0.5)*step;
5      sum+= arc(x);
6  }
```

The locally scoped variable `x` becomes private to each thread by default

```
1  !$omp parallel do
2  do i=1,num_steps
3      x=(i+0.5d0)*step
4      sum=sum+arc(x)
5  enddo
6  !$omp end parallel do
```

All variables except `i` are shared by default

parallel clauses

- ▶ Scoping clauses restricted to the lexical extent
 - ▶ `private (list)`
 - ▶ Independent variables are created for each thread
 - ▶ `shared (list)`
 - ▶ Can be used for reading and writing by multiple threads
 - ▶ `reduction (operator:variable)`
 - ▶ A private copy is made for each thread and the reduction operator is applied at the end of the parallel region
 - ▶ Operators: `+`, `*`, `-`, `&`, `|`, `^`, `&&`, `||`
 - ▶ `default (private|shared|none)`
 - ▶ By default all variables are shared
 - ▶ Locally scoped variables in C++ are private

Compute π in C++ and Fortran 90 in parallel

```
1  #pragma omp parallel for reduction(+:sum)
2  for (int i=0; i<=num_steps; i++)
3  {
4      double x=(i+0.5)*step;
5      sum+= arc(x);
6  }
```

```
1  !$omp parallel do reduction(+:sum) &
2  !$omp private(x)
3  do i=1,num_steps
4      x=(i+0.5d0)*step
5      sum=sum+arc(x)
6  enddo
7  !$omp end parallel do
```

parallel for region

```
>#edit and compile  
>qsub pi-parallel.job  
>cat pi.out  
pi is probably 3.1415926556
```

```
real    0m6.458s  
user    0m25.601s  
sys     0m0.005s
```

Loop parallelization

- ▶ Loops that can be parallelized will have the following
 - ▶ All assignments are made to arrays
 - ▶ Each element is assigned by at most one iteration
 - ▶ No iteration reads elements assigned by another iteration
 - ▶ Blocks must have one entrance and exit

Loop parallelization

- ▶ Loops with data dependence cannot be parallelized
 - ▶ $a[i] = x * a[i-1]$
 - ▶ Data dependence exists through the *dynamical* extent of a parallel region
- ▶ Loops with an undefined number of iteration cannot be parallelized
 - ▶ `while(!converged)`
- ▶ Race conditions
 - ▶ `sum=sum+i`
 - ▶ Many can be cured with the `reduction` clause
- ▶ Use control statements and divergent iterations carefully
 - ▶ `goto` must be within the block
 - ▶ `break`, `continue` must be within the block
 - ▶ `stop` and `exit` are allowed

Loop nesting

- ▶ In nested loops, both loops cannot be parallel
 - ▶ The inner loop would not be computed in parallel since all of the threads are already used in the outer loop

```
1  for (int y=0; y<25; ++y)
2  {
3      for (int x=0; x<80; ++x)
4      {
5          a[y] += b[x];
6      }
7  }
```

- ▶ In this example only the outer loop can be parallelized due to data dependence of a[y]


Exercises

Loop parallelization

- ▶ Copy `/home/sam/training/openmp/basic/`
- ▶ Parallelize the loops in `laplace2d`, `matrixvector`

Jacobi Iterations¹

- Iteratively converges to correct value (e.g. Temperature), by computing new values at each point from the average of neighboring points.



$$A_{k+1}(i,j) = \frac{A_k(i-1,j) + A_k(i+1,j) + A_k(i,j-1) + A_k(i,j+1)}{4}$$

¹Adapted from a recent PSC workshop

Loop parallelization

- ▶ Parallelize the loops in matrixvector

```
>cd matrixvector  
>#edit mxm.c  
>module purge  
>module load gcc/4.5  
>CXX mxm.c -fopenmp -o matrixvector  
>qsub matrixvector.job  
>#inspect omp.out
```